

Review

- Neural Networks
- Perceptrons
- Multilayer Perceptrons
- Project 4-b (NNets) out next week



Supervised Learning, and there's more....

- Support Vector Machines (SVMs)
- k-Nearest Neighbors (KNN)
- Ensemble methods


Support Vector Machines

- One of the most popular “off the shelf” supervised learning techniques
- Two key properties of its success
 - 1) Make non-linear problems linear, in a higher dimensional space
 - 2) Find linear separator that maximizes the margin, represent by saving all examples on the margin

The Kernel Trick



- Find a higher dimension in which the d -dimensional data is linearly separable!
- Optimal = the features that maximize the margin between the classes




SVM model: support vectors as defined by the training examples closest to the margin

Instance-based Methods

- SVM is one of a class of algorithms called instance-based
- The model is the data (as opposed to the weights in a NNet or to a DTree)


Nearest Neighbors

- Basis: a point in the data space is likely to be similarly classified as its neighbors
 - Given a new instance to classify
 - Find the points closest to it in your labeled training data, this is its “neighborhood”
 - Give it the same classification as the neighborhood



Nearest Neighbors

- Problem: how to define “neighborhood”
 - k-nearest neighbors: pick an arbitrary number of points to consider for the neighborhood
 - distance metrics: need to define what “near” means for a feature (location on a map? hair color?)



Ensemble Learning

- Why use just one hypothesis when you could use many
- EX: instead of one Decision Tree for classifying, generate several
- Save an ensemble of hypothesis and combine their predictions

Why Ensemble

- Example -- 5 classifiers instead of 1
 - classification = majority vote
 - 3 of 5 have to fail in order to misclassify
- If each h has an error $p = .1$ then the ensemble reduces this error to $p < .001$
- Works only if hypotheses are somewhat different (independent)

Boosting


- Weighted Training Set: each example has an associated weight, signifying its importance in the learning process.

Example 1	weight 1
...	...
...	...
...	...
Example i	weight i


Boosting

- Start with all $w=1$; generate h_1
- Increase weight on all examples misclassified by h_1 ; generate h_2
- Repeat to generate M hypotheses for the ensemble
- Use the ensemble by taking weighted majority classification


Boosting




Boosting



Boosting



Boosting



function ADABOOST(*examples*, L , M) **returns** a weighted-majority hypothesis

inputs: *examples*, set of N labelled examples $(x_1, y_1), \dots, (x_N, y_N)$

L , a learning algorithm

M , the number of hypotheses in the ensemble

local variables: \mathbf{w} , a vector of N example weights, initially $1/N$

\mathbf{h} , a vector of M hypotheses

\mathbf{z} , a vector of M hypothesis weights

for $m = 1$ **to** M **do**

$\mathbf{h}[m] \leftarrow L(\textit{examples}, \mathbf{w})$

$error \leftarrow 0$

for $j = 1$ **to** N **do**

if $\mathbf{h}[m](x_j) \neq y_j$ **then** $error \leftarrow error + \mathbf{w}[j]$

for $j = 1$ **to** N **do**

if $\mathbf{h}[m](x_j) = y_j$ **then** $\mathbf{w}[j] \leftarrow \mathbf{w}[j] \cdot error / (1 - error)$

$\mathbf{w} \leftarrow \text{NORMALIZE}(\mathbf{w})$

$\mathbf{z}[m] \leftarrow \log(1 - error) / error$

return WEIGHTED-MAJORITY(\mathbf{h}, \mathbf{z})

Computational Learning Theory

- How do we know if h is close to f ?
- PAC Learning (Probably Approx. Correct)
 - if h has been correct for a large number of examples it is very unlikely that it is an incorrect approximation of f

Stationary

- PAC learning assumes training and test drawn from the same distribution
- Future is like the past: therefore learned model is relevant
- Ex Computer Vision: training in different lighting than testing is non-stationary

How many examples?

- PAC learning assumes a large number of examples -- how many is needed?
 - X = all possible examples
 - D = distribution they come from
 - H = all possible hypotheses
 - N = number of training examples
- Calculate the probability that a hypothesis is bad but gets the first N samples right.
 - Depends mostly on size of H (i.e., number of features in your state)

Learning from Examples

Summary

- Supervised learning: find simple hypothesis approximately consistent with training examples
- Number examples needed is related to size of hypothesis space
- Decision Tree learning using information gain
- Neural Nets: single layer easy but limited, multilayer does non-linear, SVMs typically better
- K-nearest neighbors, training data is your model
- AdaBoost: build multiple hypotheses and vote